

Module 7: Building AI Systems

Inputs, generation, tools, evaluation, and artifacts

Module 7: Building AI Systems

Primary sources: Osmani, Chapter 7; Thomas & Hunt, Tracer Bullets review.

Learning Outcomes

- Analyze components of an AI-assisted system.
 - Describe data and context flow.
 - Evaluate integration challenges.
 - Design a basic architecture for an AI-assisted feature.
-


From Task to System

An AI-assisted system is more than a model call.

It includes:

- Input handling.
- Context assembly.
- Prompt construction.
- Model interaction.
- Tool invocation.
- Evaluation.
- Artifact tracking.
- Human oversight.

Controlled Agent Architecture



System Boundary

Define what the agent may and may not do.

For a controlled coding agent:

- Allowed tools.
 - Allowed paths.
 - Allowed commands.
 - Required approvals.
 - Required logs.
 - Required tests.
-

Context Flow

Context should be curated, not dumped.

Typical flow:

1. User request.
 2. Structured task specification.
 3. Relevant project files or summaries.
 4. Tool definitions.
 5. Evaluation criteria.
 6. Output and artifact records.
-

Integration Challenges

- Mismatched frontend/backend contracts.
 - Unclear data models.
 - Generated code that bypasses abstractions.
 - Tests that require unavailable services.
 - Tool outputs too verbose to inspect.
 - Error handling split across layers.
-

Tracer Bullet Architecture

Build one thin complete workflow first.

Example:

- Read task.
 - Ask model for plan.
 - Dispatch one safe tool.
 - Capture result.
 - Evaluate with one check.
 - Save run artifact.
-

AI System Prompting

Useful prompts ask the model to operate inside architecture.

- “Use only these tools.”
 - “Return JSON matching this schema.”
 - “If input is ambiguous, ask for clarification.”
 - “Do not modify files outside this sandbox.”
 - “Explain verification steps separately from implementation.”
-

Lab 4 Development Connection

Lab 4 begins risk analysis and production readiness.

Your risk analysis should be specific to your actual system:

- Tool misuse.
 - Sandbox escape.
 - Prompt injection.
 - False confidence from weak tests.
 - Missing human checkpoint.
-

Practice Activity

Draw your agent as boxes and arrows.

For each arrow, label:

- Data passed.
 - Trust boundary.
 - Validation step.
 - Failure mode.
-

Takeaways

- AI engineering is systems engineering.
- Context, tools, and evaluation are first-class components.
- Thin end-to-end slices make risk visible.